

Deep Learning Lab

Experiment -7

Q.1 "Image Classification Using CIFAR-10 Dataset using simple deep network with 4 hidden layers and 3 dropout layer also apply pruning and quantization to reduce size and report size of model"

1. Train the original model on CIFAR-10.
2. Save the original model (model.h5).
3. Apply pruning manually:

- If a weight is less than 0.01, we set it to 0.

Save the pruned model (pruned_model.h5).

Apply post-training quantization:

- Converts weights from 32-bit float → 8-bit int.

Save the quantized model (quantized_model.tflite).

Compare and print the sizes of all three models.

Step 1: Load CIFAR-10 Dataset

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
import os
import tempfile
import struct
```

Step 2: Normalize the Image Data

```
# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
```

```
# Normalize pixel values to [0,1]
x_train, x_test = x_train / 255.0, x_test / 255.0
# Convert labels to one-hot encoding (Fixing the issue)
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)
```

Step 3: Build and Train a Deep Neural Network

```
def create_model():
    model = keras.Sequential([
        layers.Flatten(input_shape=(32, 32, 3)), # Flatten input images
        layers.Dense(512, activation='relu'),
        layers.Dropout(0.2), # First dropout
        layers.Dense(256, activation='relu'),
        layers.Dropout(0.2), # Second dropout
        layers.Dense(128, activation='relu'),
        layers.Dense(64, activation='relu'),
        layers.Dropout(0.2), # Third dropout
        layers.Dense(10, activation='softmax') # Output layer
    ])
    # Compile the model
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model
# Create model instance
model = create_model()
# Train the model
model.fit(x_train, y_train, epochs=50, validation_data=(x_test, y_test), batch_size=64)
```

```
Epoch 1/50
782/782 ————— 9s 8ms/step - accuracy: 0.1911
- loss: 2.1619 - val_accuracy: 0.3230 - val_loss: 1.8654
Epoch 50/50
782/782 ————— 3s 4ms/step - accuracy: 0.4626
- loss: 1.5011 - val_accuracy: 0.4832 - val_loss: 1.4592
```

Step 4: Save the Model

```
_, model_file = tempfile.mkstemp('.h5') # Create temporary file
model.save(model_file) # Save model
original_size = os.path.getsize(model_file) / (1024 * 1024) # Convert to MB
print(f"Original Model Size: {original_size:.2f} MB")
```

OUTPUT

Original Model Size: 20.03 MB

Step 5: Apply Model Pruning (Reducing Unimportant Weights)

```
pruned_model = tf.keras.models.clone_model(model)
pruned_model.set_weights([np.where(np.abs(w) > 0.01, w, 0) for w in
model.get_weights()])
pruned_model.save("pruned_model.h5")
```

Step 6: Apply Quantization (Reducing Precision of Weights)

```
# Apply quantization (convert model to TensorFlow Lite format)
converter = tf.lite.TFLiteConverter.from_keras_model(pruned_model)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
quantized_model = converter.convert()
# Save quantized model
with open("quantized_model.tflite", "wb") as f:
    f.write(quantized_model)

original_size = os.path.getsize("model.h5") / 1024 # Convert to KB
pruned_size = os.path.getsize("pruned_model.h5") / 1024
quantized_size = os.path.getsize("quantized_model.tflite") / 1024

print(f"Original Model Size: {original_size:.2f} KB")
print(f"Pruned Model Size: {pruned_size:.2f} KB")
print(f"Quantized Model Size: {quantized_size:.2f} KB")
```

Step 7: Compare Model Size

Original Model Size: 20512.41 KB

Pruned Model Size: 6856.16 KB

Quantized Model Size: 1724.73 KB

Pruning: Pruning removes unnecessary weights (connections) in a neural network by setting small weights to zero. This reduces model size and speeds up inference while maintaining accuracy. It helps in optimizing storage and computational efficiency.

Quantization: Quantization reduces the precision of model weights and activations (e.g., from 32-bit floating-point to 8-bit integers). This significantly decreases the model size and makes it faster, especially for deployment on mobile and edge devices.

Thank You Sir

